

Department of Computer Science
Database and Information Systems Group

Project Thesis:

**Performance Evaluation of Different Open
Source Implementations of Data Structures
and Other Algorithms in the context of a
DBMS Buffer Manager**

by **Max Fabian Gilbert***

Day of release: January 31, 2019

Abstract

Needless to say, every database management system needs to be able to manage data. The data structures used to manage those data in a database have a major influence on various characteristics (e.g. performance) of a database management system and therefore, the usage of specific data structures (e.g. B-tree indexes) and even some implementation details of those are very important decisions in DBMS design.

But for correct and performant operation, a DBMS needs to manage various kinds of meta data as well. Some of those meta data needs to be persistent (e.g. the catalog of a relational DBMS) but some can also be non-persistent. Because of the non-persistence of data managed by the buffer management of a DB, the meta data required for the buffer manager are also usually non-persistent. The data structures used to manage those meta data are—unlike the data structures used to manage the data—more an implementation than a design decision. For some kinds of those meta data, it's—due to the non-criticality of the specific meta data management—even reasonable to use data structures provided by the used programming language even though there might be more performant data structures for the purpose. But more performant implementations for most of those data structures don't need to be implemented specifically for one project, there are many different implementations available in open source and proprietary libraries.

This work is a performance evaluation of various MPMC

This page intentionally left blank.

Contents

1	Buffer Frame Free List	1
1.1	Purpose	1
1.2	Compared Queue Implementations	1
1.2.1	Boost Lock-Free Queue with variable size	2
1.2.2	Boost Lock-Free Queue with fixed size	2
1.2.3	CDS Basket Lock-Free Queue	3
1.2.4	CDS Flat-Combining Lock-Free Queue	3
1.2.5	CDS Michael & Scott Lock-Free Queue	4
1.2.6	CDS Variation of Michael & Scott Lock-Free Queue	4
1.2.7	CDS Michael & Scott Blocking Queue with Fine-Grained Locking	5
1.2.8	CDS Ladan-Mozes & Shavit Optimistic Queue	5
1.2.9	CDS Segmented Queue	5
1.2.10	CDS Vyukov's MPMC Bounded Queue	6
1.2.11	Folly MPMC Queue	6
1.2.12	Dmitry Vyukov's Bounded MPMC Queue	7
1.2.13	Gavin Lambert's MPMC Bounded Lock-Free Queue	7
1.2.14	moodycamel::ConcurrentQueue	7
1.2.15	Matt Stump's Bounded MPMC Queue	8
1.2.16	Erik Rigtorp's Bounded MPMC Queue	8
1.2.17	Threading Building Blocks Concurrent Queue	8
1.2.18	Threading Building Blocks Bounded Concurrent Dual Queue	8
1.3	Performance Evaluation	9
1.3.1	Micro Benchmark	9
1.3.2	Used Versions of the Libraries and Queue Implementations	9
1.3.3	Configuration of the Used System	10
1.4	Conclusion	10

2	RANDOM Page Eviction	11
2.1	Purpose	11
2.2	Compared Pseudorandom Number Generators	11
2.2.1	std::minstd_rand0	12
2.2.2	std::minstd_rand	12
2.2.3	std::mt19937	12
2.2.4	std::mt19937_64	12
2.2.5	std::ranlux24_base	12
2.2.6	std::ranlux48_base	13
2.2.7	std::ranlux24	13
2.2.8	std::ranlux48	13
2.2.9	std::knuth_b	13
2.2.10	std::rand	13
2.2.11	std::random_device	13
2.2.12	Xorshift32	13
2.2.13	Xorshift64	14
2.2.14	Xorshift96	14
2.2.15	Xorshift128	15
2.2.16	XorWow	16
2.2.17	Xorshift*	16
2.2.18	Xorshift+	16
2.2.19	Xoroshiro128+	16
2.3	Performance Evaluation	16
2.3.1	Micro Benchmark	16
2.3.2	Configuration of the Used System	16
2.4	Conclusion	16
3	LOOP Page Eviction	17
3.1	Purpose	17
3.2	Compared Concurrent Counter Implementations	17
3.2.1	Mutex Counter	17
3.2.2	Spinlock Counter	17
3.2.3	Modulo Counter	18
3.2.4	Local Counter	18
3.2.5	Local Modulo Counter	18
3.2.6	Clunky Counter	19

Contents

3.3	Performance Evaluation	19
3.3.1	Microbenchmark	19
3.3.2	Configuration of the Used System	20
3.4	Conclusion	20
	Bibliography	21

This page intentionally left blank.

1 Buffer Frame Free List

1.1 Purpose

Every disk-based DBMS requires some kind of buffer manager. A disk-based DBMS stores the pages of a database on secondary storage but to read and write pages, they are required to be in memory.

This feature is provided by the buffer pool management by managing the currently used subset of the database pages in buffer frames located in memory. A buffer frame is a portion of memory that can hold one database page and each of those frames got a frame index as identifier.

During operation, database pages are dynamically fetched from the database into buffer frames. Once a page is not required anymore, it might be evicted from the buffer pool freeing a buffer frame.

Due to the fact that pages are only allowed to be fetched into free buffer frames, the buffer manager needs to know all the free buffer frames. Therefore, a free list for the buffer frames—storing the frame indexes of free buffer frames—is required.

1.2 Compared Queue Implementations

To ease implementation of page eviction strategies like CLOCK, a free list should use a FIFO data structure like a queue. Therefore the buffer frame freed first is (re-)used first as well.

Almost every state-of-the-art DBMS support multithreading and therefore, there are usually multiple threads concurrently fetching pages into the buffer pool and evicting pages from the buffer pool. Following this, a buffer frame free list has to support thread-safe functions to push frame indexes to the free list and to pop frame indexes from it. Queues providing those thread-safe access functions are usually called multi-producer (add

frame indexes) multi-consumer (retrieve/remove frame indexes) queues (MPMC queues).

An approximate number of buffer indexes in the free list must also be provided by any free list implementation to support the eviction of pages once there are only a few free buffer frames left. Thread-safe access to this number is desirable but not absolutely required.

1.2.1 Boost Lock-Free Queue with variable size

The famous *Boost C++ Libraries*¹ offer a lock-free unbounded MPMC queue² in the library `Boost.Lockfree`³. Like many other non-blocking thread-safe data structures, this MPMC queue uses atomic operations instead of locks or mutexes. To support queues of dynamically changing sizes, this queue implementation also uses a free list for the dynamic memory management internally.

This data structure does not offer the number of contained elements and therefore, an approximate number of buffer indexes in the free list needs to be managed outside.

1.2.2 Boost Lock-Free Queue with fixed size

This data structure is identical to the data structure in Subsection 1.2.1 but does not use dynamic memory management internally—it is a bounded queue. Therefore, the capacity of the queue (i.e. the maximum number of buffer frames of the buffer pool) needs to be specified beforehand which allows the usage of a fixed-size array instead of dynamically allocated nodes.

¹<https://www.boost.org/>

²<https://www.boost.org/doc/libs/release/doc/html/boost/lockfree/queue.html>

³<https://www.boost.org/doc/libs/release/doc/html/lockfree.html>

1.2.3 CDS Basket Lock-Free Queue

Besides other concurrent data structures, the *Concurrent Data Structures* C++ library⁴ offers many different thread-safe queue implementations. The unbounded *Basket Lock-Free Queue*⁵ is based on the algorithm proposed by M. Hoffman, O. Shalev and N. Shavit in [HSS07].

Internally, this queue does not use an absolute FIFO order. Instead, it puts concurrently enqueued elements into one “basket” of elements. The elements within one basket are not specifically ordered but the different “baskets” used over time are ordered according to FIFO. Therefore, the dequeue operation just dequeues one of the elements in the oldest “basket”. The dynamic memory management uses a garbage collector to deallocate emptied “baskets”.

1.2.4 CDS Flat-Combining Lock-Free Queue

The *Concurrent Data Structures* C++ library does also offer an unbounded thread-safe queue that uses Flat Combining⁶. The Flat Combining technique was proposed by D. Hendler, I. Incze, N. Shavit and M. Tzafrir in [Hen+10]. This technique is used to make any sequential data structure thread-safe—in case of the *Flat-Combining Lock-Free Queue*, the `std::queue`⁷ of the *C++ Standard Library*⁸ is used as base data structure.

The Flat Combining technique uses thread-local publication lists to record operations performed by those threads. A global lock is needed to be acquired to combine these thread-local publication lists into the global, sequential data structure. The thread which acquired the global lock also combines the publication lists of all other threads reducing the locking overhead. The returned value of each operation executed during the combining is stored into the respective publication list together with the global combining pass number. A thread with a non-empty publication list that

⁴<https://github.com/khizmax/libcds>

⁵http://libcds.sourceforge.net/doc/cds-api/classcds_1_1container_1_1_basket_queue.html

⁶http://libcds.sourceforge.net/doc/cds-api/classcds_1_1container_1_1_f_c_queue.html

⁷<https://en.cppreference.com/w/cpp/container/queue>

⁸<https://en.cppreference.com/w/cpp>

cannot acquire the global lock needs to wait till the combining thread updated its publication list.

1.2.5 CDS Michael & Scott Lock-Free Queue

Another unbounded lock-free queue implementation offered by the *Concurrent Data Structures* C++ library is based on the famous Michael & Scott lock-free queue algorithm⁹ which was proposed by M. Michael and M. Scott in [MS96].

The Michael & Scott lock-free queue basically uses compare-and-swap (CAS) operations on the tail of the queue to synchronize enqueue operations. If a thread reads a NULL value as next element after the queue's tail, it swaps this value atomically with the value enqueued by this thread. Afterwards it adjusts the tail pointer. If a thread does not read the NULL value there during the CAS operation, another thread has not already adjusted the tail pointer and this thread needs to retry its enqueue operation with the new tail pointer. The dequeue operation is implemented similarly. The memory occupied by already dequeued elements is deallocated using a garbage collector provided by the library.

1.2.6 CDS Variation of Michael & Scott Lock-Free Queue

The *Concurrent Data Structures* C++ library also offers an optimized variation of the Michael & Scott unbounded lock-free queue algorithm¹⁰ which is based on the works of S. Doherty, L. Groves, V. Luchangco and M. Moir in [Doh+04].

This optimization of the Michael & Scott lock-free queue optimizes the dequeue operation to only read the tail pointer once.

⁹http://libcdfs.sourceforge.net/doc/cds-api/classcds_1_1container_1_1_m_s_queue.html

¹⁰http://libcdfs.sourceforge.net/doc/cds-api/classcds_1_1container_1_1_moir_queue.html

1.2.7 CDS Michael & Scott Blocking Queue with Fine-Grained Locking

M. Michael and M. Scott did also propose a blocking queue algorithm in [MS96]. This unbounded blocking queue implementation¹¹ is also offered by the *Concurrent Data Structures* C++ library.

This blocking queue algorithm uses one read and one write lock protecting the head and tail of the queue. Therefore, only one thread at a time can enqueue and only one thread at a time can dequeue elements. The deallocation of memory during dequeuing is done by the dequeuing thread instead of relying on a garbage collector.

1.2.8 CDS Ladan-Mozes & Shavit Optimistic Queue

The *Concurrent Data Structures* C++ library also offers an unbounded optimistic queue implementation¹² which is based on an algorithm proposed by E. Ladan-Mozes and N. Shavit in [LS04].

Instead of using expensive CAS operations on a singly-linked list (like in the Michael & Scott lock-free queue), this algorithm uses a doubly-linked list with the possibility to detect and fix inconsistent enqueue and dequeue operations. Deallocation of memory is done using a garbage collector.

1.2.9 CDS Segmented Queue

The unbounded segmented queue implementation¹³ of the *Concurrent Data Structures* C++ library is based on an algorithm proposed by Y. Afek, G. Korland and E. Yanovsky in [AKY10].

This thread-safe queue algorithm is very similar to the basket lock-free queue. It also uses a relaxed FIFO order by ordering segments containing multiple elements instead of single elements. A thread enqueueing or

¹¹http://libcdfs.sourceforge.net/doc/cds-api/classcds_1_1container_1_1_r_w_queue.html

¹²http://libcdfs.sourceforge.net/doc/cds-api/classcds_1_1container_1_1_optimistic_queue.html

¹³http://libcdfs.sourceforge.net/doc/cds-api/classcds_1_1container_1_1_segmented_queue.html

dequeuing elements into the tail segment or from the head segment selects one of the slots inside the segment randomly. CAS operations are used to atomically enqueue or dequeue an element from a slot. If the CAS fails, another slot is taken randomly. The size of each segment—which can be selected (8 was used for the performance evaluation in Section 1.3)—determines the relaxiness of the FIFO order. Deallocation of emptied segments is done by a garbage collector.

1.2.10 CDS Vyukov’s MPMC Bounded Queue

The last thread-safe queue implementation¹⁴ provided by the *Concurrent Data Structures* C++ library is bounded and was developed by D. Vyukov¹⁵. The queue implementation in Subsection 1.2.12 is his original implementation.

Vyukov’s thread-safe queue implementation is very similar to Michael & Scott blocking queue with fine-grained locking from Subsection 1.2.7 but instead of using mutexes as locks, his implementation uses atomic read-modify-write (**RMW**) operations. This results in a cost of basically one CAS operation per enqueue/dequeue operation.

1.2.11 Folly MPMC Queue

Facebook’s open source library *Folly*¹⁶ provides a bounded lock-free queue implementation. An unbounded queue is also provided but due to the typically lower performance of unbounded queues, it is not evaluated in Section 1.3.

Folly’s MPMC queue uses a ticket dispenser system to give a thread access to one of the single-element queues used. Those ticket dispensers for the head and tail of the queue use atomic increment operations which are supposed to be more robust to contention than CAS operations used e.g. in the Michael & Scott lock-free queue.

¹⁴http://libcdis.sourceforge.net/doc/cds-api/classcds_1_1container_1_1_vyukov_m_p_m_c_cycle_queue.html

¹⁵<http://www.1024cores.net/home/lock-free-algorithms/queues/bounded-mpmc-queue>

¹⁶<https://github.com/facebook/folly>

1.2 Compared Queue Implementations

1.2.12 Dmitry Vyukov's Bounded MPMC Queue

This¹⁷ is Vyukov's original implementation of his bounded thread-safe MPMC queue.

1.2.13 Gavin Lambert's MPMC Bounded Lock-Free Queue

This¹⁸ is another implementation of Vyukov's thread-safe queue made by [Gavin Lambert.

1.2.14 `moodycamel::ConcurrentQueue`

This lock-free queue implementation¹⁹ is either unbounded or bounded depending on the used enqueueing functions and on the optional preallocation of memory (bounded behavior is used during the performance evaluation in Section 1.3). A blocking queue implementation provided by the same library is not evaluated in Section 1.3 because it is just a wrapper around the non-blocking version adding additional overhead in low-contention workloads (like the free list).

Internally, this queue implementation uses one SPMC (single producer/multiple consumer) queue per thread. Each thread enqueues elements only into its thread-local SPMC queue. When a thread tries to dequeue an element, it checks SPMC queues for emptiness until it finds one containing elements. It then dequeues one element from the SPMC queue. Therefore, this thread-safe queue does not maintain the order of elements enqueued by different threads.

Due to the implementation using multiple SPMC queues, this queue implementation should only be used as a buffer frame free list when there is exactly one thread evicting pages from the buffer pool—and therefore, enqueueing buffer frame indexes of emptied buffer frames.

¹⁷<http://www.1024cores.net/home/lock-free-algorithms/queues/bounded-mpmc-queue>

¹⁸<https://gist.github.com/uecasmb547db812ae4bba39bb1bd0443801507>

¹⁹<https://github.com/cameron314/concurrentqueue>

1.2.15 Matt Stump's Bounded MPMC Queue

This²⁰ is another implementation of Vyukov's thread-safe queue made by Matt Stump.

1.2.16 Erik Rigtorp's Bounded MPMC Queue

The bounded lock-free queue²¹ of Erik Rigtorp uses a ticket dispenser system similar to the one of *Folly*'s MPMC queue from Subsection 1.2.11.

1.2.17 Threading Building Blocks Concurrent Queue

The *Threading Building Blocks* library²² is an open source library originally developed by Intel®. The first thread-safe queue implementation²³ of this library is unbounded and non-blocking.

Internally, this queue implementation uses multiple lock-based micro queues to allow concurrent enqueue/dequeue executions. Therefore, the guarantees of this queue is similar to those of the `moodycamel::ConcurrentQueue` from Subsection 1.2.14.

1.2.18 Threading Building Blocks Bounded Concurrent Dual Queue

The other thread-safe queue implementation²⁴ of the *Threading Building Blocks* library is unbounded and partially non-blocking.

This queue implementation is almost identical to the other one of the *Threading Building Blocks* library but it does allow the limitation of the capacity. An enqueueing operation has to wait if the queue is already full according to the specified capacity.

²⁰<https://github.com/mstump/queues>

²¹<https://github.com/rigtorp/MPMCQueue>

²²<https://www.threadingbuildingblocks.org/>

²³<https://software.intel.com/en-us/node/506200>

²⁴<https://software.intel.com/en-us/node/506201>

1.3 Performance Evaluation

1.3.1 Micro Benchmark

The used micro benchmark simulates a high contented free list. The number of working threads, the number of iterations (either the fetching of a page into a free buffer frame or the eviction of a batch of pages) per thread and the batch size of buffer frames to be freed at once can be varied. It does not simulate a complete buffer pool—there is only the free list with operations to enqueue and dequeue buffer frame indexes. Each working thread performs the following operations per iterations:

- If the free list is not empty:
 - Retrieve a buffer frame index from the free list.
 - Mark the retrieved buffer frame used.
- If the free list is empty:
 - While the free list is smaller than the batch eviction size:
 - * Select a random buffer frame index using a fast random numbers generator.
 - * If this buffer frame index is marked used:
 - Mark the selected buffer frame index unused.
 - Add the selected buffer frame index to the free list.

1.3.2 Used Versions of the Libraries and Queue Implementations

- *Boost C++ Libraries* 1.58
- *Concurrent Data Structures* C++ library 2.3.3²⁵
- *Folly* a15fcb1e76²⁶
- Dmitry Vyukov's Original MPMC Queue as of September 2017

²⁵<https://github.com/khizmax/libcds/tree/5fc87a172bd82f8a7040b8b83f32ce0e635e82ea>

²⁶<https://github.com/facebook/folly/tree/a15fcb1e76444f7d464b263ad37bf3b5fbfdf33e>

- Gavin Lambert's MPMC Queue as of September 2017²⁷
- moodycamel::ConcurrentQueue 9f9c4e0cf4²⁸
- Matt Stump's MPMC Queue 319c253d68²⁹
- Erik Rigtorp's MPMC Queue 57366e41f3³⁰
- Intel® Threading Building Blocks 2017 Update 7

1.3.3 Configuration of the Used System

- **CPU:** $2 \times$ Intel® Xeon® Processor X5670 @ $6 \times 2.93\text{GHz}$ released early 2010
- **Main Memory:** $12 \times 8\text{GB} = 96\text{GB}$ of DDR2-SDRAM @1333MHz
- **OS:** Ubuntu 16.04

1.4 Conclusion

²⁷<https://gist.github.com/uecasmb547db812ae4bba39bb1bd0443801507/e40906811cb14118d328c353250559fe359f3ba7>

²⁸<https://github.com/cameron314/concurrentqueue/tree/9f9c4e0cf400bcc5c27a041e524f04e950736b25>

²⁹<https://github.com/mstump/queues/tree/319c253d68f14ac9593c3727d1597a87af73c99b>

³⁰<https://github.com/rigtorp/MPMCQueue/tree/57366e41f3f48316f175c2e704795f519a92e1d5>

2 RANDOM Page Eviction

2.1 Purpose

The buffer manager of a DBMS needs to evict pages from buffer frames when currently not buffered pages need to be fetched from the database while there are no more free buffer frames. Every buffer manager got a page evictionner—implementing one of the many page eviction algorithms developed since the 1960s—for that purpose.

According to Belady's classification in [Bel66], the RANDOM eviction algorithm is the most representative algorithm in his *Class 1* of page eviction algorithms. Those *Class 1* page eviction algorithms do not use any information about the usage of a buffered page but just apply a static rule for the eviction decision. According to the newer classification of Effelsberg and Härder in [EH84], the RANDOM eviction algorithm is the only algorithm in the class of algorithms using neither the age of a buffered page nor the references of it for the eviction decision.

The RANDOM strategy is the simplest page eviction strategy possible resulting in a low overhead and bad hit rates.

2.2 Compared Pseudorandom Number Generators

The only operation performed by the RANDOM page evictionner to decide which page to eviction from the buffer pool is the generation of a pseudorandom number in the range of buffer frame indexes. The database page contained in the selected buffer frame is evicted afterwards.

There are many different classes of pseudorandom number generators (**PRNG**). Some of them provide pseudorandom numbers of high randomness others just take only few CPU cycles and almost no memory to generate

a random number.

2.2.1 `std::minstd_rand0`

The *C++ Standard Library*¹ provides the original *MINSTD* PRNG as proposed by S. Park and K. Miller in [PM88] predefined in the header `<random>`².

The original *MINSTD* is a *Lehmer random number generator* which is a linear congruential generator (LCG). It uses the parameters $m = 2^31 - 1$, $a = 16807$ and $c = 0$.

2.2.2 `std::minstd_rand`

The *C++ Standard Library*¹ also provides the revised *MINSTD* PRNG as proposed by S. Park, K. Miller and P. Stockmeyer in [PMS93] predefined in the header `<random>`².

This revised *MINSTD* improves the quality of generated pseudorandom numbers by using the parameter $a = 48271$.

2.2.3 `std::mt19937`

The standard implementation of the Mersenne Twister MT19937—proposed by M. Matsumoto and T. Nishimura in [MN98]—is also provided in the header `<random>`² of the *C++ Standard Library*¹.

2.2.4 `std::mt19937_64`

The 64-bit version of the Mersenne Twister—MT19937-64—is also provided in the header `<random>`² of the *C++ Standard Library*¹.

2.2.5 `std::ranlux24_base`

This PRNG, based on the 24-bit RANLUX generator—proposed in [Lüs94] by M. Lüscher—, is a *Subtract-With-Borrow* generator provided by the *C++ Standard Library*¹ in header `<random>`².

¹<https://en.cppreference.com/w/cpp>

²<https://en.cppreference.com/w/cpp/numeric/random>

2.2 Compared Pseudorandom Number Generators

The computational cost of this PRNG is very high but the pseudorandom numbers generated are of high randomness.

2.2.6 `std::ranlux48_base`

This PRNG is like the PNRG from Subsection 2.2.5 but it uses 48-bit instead of 24-bit words. It is provided by the *C++ Standard Library*¹ in header `<random>`².

2.2.7 `std::ranlux24`

This is the 24-bit RANLUX generator originally proposed by M. Lüscher and F. James in [Lüs94] and [Jam94] provided by the *C++ Standard Library*¹ in header `<random>`².

2.2.8 `std::ranlux48`

This is the 48-bit version of the PRNG from Subsection 2.2.7 provided by the *C++ Standard Library*¹ in header `<random>`².

2.2.9 `std::knuth_b`

2.2.10 `std::rand`

2.2.11 `std::random_device`

2.2.12 Xorshift32

```
1  thread_local bool seed_initialized;
2  thread_local uint32_t seed;
3
4  uint32_t xorshift32_random() {
5      if (!seed_initialized) {
6          seed = std::random_device{}();
7          seed_initialized = true;
8      }
9      seed ^= seed << 13;
10     seed ^= seed >> 17;
11     seed ^= seed << 5;
12     return (seed % (block_count - 1) + 1);
13 }
```

2.2.13 Xorshift64

```

1  thread_local bool seed_initialized;
2  thread_local uint64_t seed;

4  uint32_t xorshift64_random() {
5      if (!seed_initialized) {
6          seed = std::random_device{}();
7          seed_initialized = true;
8      }
9      seed ^= seed << 13;
10     seed ^= seed >> 7;
11     seed ^= seed << 17;
12     return (seed % (block_count - 1) + 1);
13 }

```

2.2.14 Xorshift96

```

1  thread_local bool seed_initialized;
2  thread_local uint32_t seed_0;
3  thread_local uint32_t seed_1;
4  thread_local uint32_t seed_2;

6  uint32_t xorshift96_random() {
7      if (!seed_initialized) {
8          seed_0 = std::random_device{}();
9          seed_1 = std::random_device{}();
10         seed_2 = std::random_device{}();
11         seed_initialized = true;
12     }
13     uint32_t t;
14     seed_0 ^= seed_0 << 16;
15     seed_0 ^= seed_0 >> 5;
16     seed_0 ^= seed_0 << 1;

18     t = seed_0;
19     seed_0 = seed_1;
20     seed_1 = seed_2;
21     seed_2 = t ^ seed_0 ^ seed_1;

23     return (seed_2 % (block_count - 1)) + 1;

```

2.2 Compared Pseudorandom Number Generators

24 }

2.2.15 Xorshift128

```
1  thread_local bool seed_initialized;
2  thread_local uint32_t seed_0;
3  thread_local uint32_t seed_1;
4  thread_local uint32_t seed_2;
5  thread_local uint32_t seed_3;
6
7  uint32_t xorshift128_random() {
8      if (!seed_initialized) {
9          seed_0 = std::random_device{}();
10         seed_1 = std::random_device{}();
11         seed_2 = std::random_device{}();
12         seed_3 = std::random_device{}();
13         seed_initialized = true;
14     }
15     uint32_t t = seed_0 ^ (seed_0 << 11);
16     seed_0 = seed_1;
17     seed_1 = seed_2;
18     seed_2 = seed_3;
19
20     seed_3 ^= (seed_3 >> 19) ^ t ^ (t >> 8);
21
22     return (seed_3 % (block_count - 1) + 1);
23 }
```

2.2.16 XorWow

2.2.17 Xorshift*

2.2.18 Xorshift+

2.2.19 Xoroshiro128+

2.3 Performance Evaluation

2.3.1 Micro Benchmark

2.3.2 Configuration of the Used System

- **CPU:** $2 \times$ Intel® Xeon® Processor X5670 @ 6×2.93 GHz released early 2010
- **Main Memory:** 12×8 GB = 96GB of DDR2-SDRAM @1333MHz
- **OS:** *Ubuntu 16.04*

2.4 Conclusion

3 LOOP Page Eviction

3.1 Purpose

The LOOP eviction algorithm is the simplest form of RANDOM eviction algorithms where the eviction candidates are selected round-robin according to the buffer frame index.

The LOOP eviction strategy uses a “pseudorandom” number generator—basically a counter counting modulo the number of buffer pool frames—which generates an ordered sequence of buffer indexes.

3.2 Compared Concurrent Counter Implementations

3.2.1 Mutex Counter

```
1 inline uint32_t mutex_counter() {  
2     static std::mutex idx_lock;  
3     static uint32_t used_idx = 0;  
4     std::lock_guard<std::mutex> guard(idx_lock);  
5     used_idx++;  
6     if (used_idx >= block_cnt) {  
7         used_idx = 1;  
8     }  
9     return used_idx;  
10 }
```

3.2.2 Spinlock Counter

```
1 inline uint32_t spinlock_counter() {  
2     static std::atomic_flag idx_lock = ATOMIC_FLAG_INIT;  
3     static uint32_t used_idx = 0;  
4     uint32_t this_idx;  
5 }
```



```

5  while (idx_lock.test_and_set(std::memory_order_acquire)) {}
6  used_idx++;
7  if (used_idx >= block_cnt) {
8      used_idx = 1;
9  }
10 this_idx = used_idx;
11 idx_lock.clear(std::memory_order_release);
12 return this_idx;
13 }

```

3.2.3 Modulo Counter

```

1  inline uint32_t modulo_counter() {
2      static std::atomic<uint64_t> used_idx;
3      while (true) {
4          uint32_t this_idx = used_idx++ % block_cnt;
5          if (this_idx != 0) {
6              return this_idx;
7          } else {
8              continue;
9          }
10     }
11 }

```

3.2.4 Local Counter

```

1  inline uint32_t local_counter() {
2      static thread_local uint32_t used_idx = 0;
3      used_idx++;
4      if (used_idx >= block_cnt) {
5          used_idx = 1;
6      }
7      return used_idx;
8  }

```

3.2.5 Local Modulo Counter

```

1  inline uint32_t local_modulo_counter() {
2      static thread_local uint64_t used_idx = 0;
3      static uint32_t modulo_factor = block_cnt - 1;

```

3.3 Performance Evaluation

```
4     return (used_idx++ % modulo_factor) + 1;  
5 }
```

3.2.6 Clunky Counter

```
1 inline uint32_t clunky_counter() {  
2     static std::atomic<uint32_t> used_idx(1);  
3     uint32_t picked_idx = used_idx;  
4     if (picked_idx < block_cnt) {  
5         used_idx++;  
6         return picked_idx;  
7     } else {  
8         return used_idx = 1;  
9     }  
10 }
```

3.3 Performance Evaluation

Due to the implementation independence of the eviction decisions of LOOP page evictioners, the achieved hit rates achieved in an exemplary DBMS are not required as performance measure for this evaluation. The only performance difference between the different LOOP page evictioners is the overhead imposed by the concurrent counting. Therefore, a microbenchmark measuring only the execution time of the counting is appropriate.

The variables of the evaluation are the concurrent counter implementation—the alternatives presented in the previous section are evaluated—and the number of threads concurrently incrementing the counter. The smallest (> 0) and largest integers returned by the counters—representing the smallest and largest buffer pool indexes—do not significantly influence the performance of the LOOP page evictioners. Therefore, this integer interval $[1..999]$ is a constant in this evaluation.

3.3.1 Microbenchmark

The microbenchmark used for the performance evaluation of the LOOP page eviction algorithms forks a given number of worker threads—each

of them calls the evaluated concurrent counter a given number of times (1,000,000)—and measures the wall time elapsed until all the worker threads finished their operation and joined.

3.3.2 Configuration of the Used System

- **CPU:** $2 \times$ Intel® Xeon® Processor X5670 @ $6 \times 2.93\text{GHz}$ released early 2010
- **Main Memory:** $12 \times 8\text{GB} = 96\text{GB}$ of DDR2-SDRAM @1333MHz
- **OS:** *Ubuntu 18.04*

3.4 Conclusion

Bibliography

- [AKY10] Yehuda Afek, Guy Korland, and Eitan Yanovsky. “Quasi-Linearizability: Relaxed Consistency for Improved Concurrency”. In: *Principles of Distributed Systems*. Lecture Notes in Computer Science (6490 2010): *14th International Conference, OPODIS 2010 Tozeur, Tunisia, December 14-17, 2010 Proceedings*. Ed. by Chenyang Lu, Toshimitsu Masuzawa, and Mohamed Mosbah. Found. by Gerhard Goos, Juris Hartmanis, and Jan van Leeuwen, pp. 395–410. DOI: 10.1007/978-3-642-17653-1_29. URL: <http://mcg.cs.tau.ac.il/papers/opodis2010-quasi.pdf> (visited on Jan. 24, 2019).
- [Bel66] Laszlo A. Belady. “A Study of Replacement Algorithms for Virtual-Storage Computer”. In: *IBM Systems Journal* 5 (2 June 1966), pp. 78–101. ISSN: 0018-8670. DOI: 10.1147/sj.52.0078.
- [Doh+04] Simon Doherty et al. “Formal Verification of a Practical Lock-Free Queue Algorithm”. In: *Formal Techniques for Networked and Distributed Systems – FORTE 2004*. Lecture Notes in Computer Science (3235 2004): *24th IFIP WG 6.1 International Conference, Madrid Spain, September 27-30, 2004 Proceedings*. Ed. by David de Frutos-Escrig and Manuel Núñez. Found. by Gerhard Goos, Juris Hartmanis, and Jan van Leeuwen, pp. 97–114. DOI: 10.1007/978-3-540-30232-2_7.
- [EH84] Wolfgang Effelsberg and Theo Härder. “Principles of Database Buffer Management”. In: *ACM Transactions on Database Systems* 9 (4 Dec. 1984). Ed. by Robert W. Taylor, pp. 560–595. ISSN: 0362-5915. DOI: 10.1145/1994.2022. URL: <http://dl.acm.org/citation.cfm?id=2022> (visited on Feb. 2, 2017).

- [Hen+10] Danny Hendler et al. “Flat Combining and the Synchronization-Parallelism Tradeoff”. In: (2010): *SPAA '10: Proceedings of the Twenty-second Annual ACM Symposium on Parallelism in Algorithms and Architectures*. Ed. by Friedhelm Meyer auf der Heide and Cynthia Phillips, pp. 355–364. doi: 10.1145/1810479.1810540. URL: <https://www.cs.bgu.ac.il/~hendlerd/papers/flat-combining.pdf> (visited on Jan. 24, 2019).
- [HSS07] Moshe Hoffman, Ori Shalev, and Nir Shavit. “The Baskets Queue”. In: *Principles of Distributed Systems*. Lecture Notes in Computer Science (4878 2007): *11th International Conference, OPODIS 2007 Guadeloupe, French West Indies, December 17-20, 2007 Proceedings*. Ed. by Eduardo Tovar, Philippas Tsigas, and Hacène Fouchal. Found. by Gerhard Goos, Juris Hartmanis, and Jan van Leeuwen, pp. 401–414. doi: 10.1007/978-3-540-77096-1_29. URL: <https://people.csail.mit.edu/shanir/publications/Baskets%20Queue.pdf> (visited on Jan. 24, 2019).
- [Jam94] Frederick E. James. “RANLUX: A Fortran Implementation of the High-Quality Pseudorandom Number Generator of Lüscher”. In: *Computer Physics Communications* 79 (1 Feb. 1994), pp. 111–114. ISSN: 0010-4655. doi: 10.1016/0010-4655(94)90233-X.
- [LS04] Edya Ladan-Mozes and Nir Shavit. “An Optimistic Approach to Lock-Free FIFO Queues”. In: *Distributed Computing*. Lecture Notes in Computer Science (3274 2004): *18th International Conference, DISC 2004, Amsterdam, The Netherlands, October 4-7, 2004 Proceedings*. Ed. by Rachid Guerraoui. Found. by Gerhard Goos, Juris Hartmanis, and Jan van Leeuwen, pp. 117–131. doi: 10.1007/978-3-540-30186-8_9. URL: http://people.csail.mit.edu/shanir/publications/FIFO_Queues.pdf (visited on Jan. 24, 2019).
- [Lüs94] Martin Lüscher. “A Portable High-Quality Random Number Generator for Lattice Field Theory Simulations”. In: *Computer*

Bibliography

- Physics Communications* 79 (1 Feb. 1994), pp. 100–110. ISSN: 0010-4655. DOI: 10.1016/0010-4655(94)90232-1.
- [MN98] Makoto Matsumoto and Takuji Nishimura. “Mersenne Twister: A 623-Dimensionally Equidistributed Uniform Pseudo-Random Number Generator”. In: *ACM Transactions on Modeling and Computer Simulation (TOMACS)* 8 (1 Jan. 1998): *Special Issue on Uniform Random Number Generation*. Ed. by Philip Heidelberger, Raymond Coutre, and Pierre L’Ecuyer, pp. 3–30. ISSN: 1049-3301. DOI: 10.1145/272991.272995. URL: <http://www.math.sci.hiroshima-u.ac.jp/~m-mat/MT/ARTICLES/mt.pdf> (visited on Jan. 24, 2019).
- [MS96] Maged M. Michael and Michael L. Scott. “Simple, Fast, and Practical Non-Blocking and Blocking Concurrent Queue Algorithms”. In: (1996): *PODC ’96: Proceedings of the Fifteenth Annual ACM Symposium on Principles of Distributed Computing*. Ed. by James E. Burns and Yoram Moses, pp. 267–275. DOI: 10.1145/248052.24810629. URL: http://www.cs.rochester.edu/~scott/papers/1996_PODC_queues.pdf (visited on Jan. 24, 2019).
- [PM88] Stephen K. Park and Keith W. Miller. “Random Number Generators: Good Ones Are Hard To Find”. In: *Communications of the ACM* 31 (10 Oct. 1988), pp. 1192–1201. ISSN: 0001-0782. DOI: 10.1145/63039.63042. URL: <http://www.firstpr.com.au/dsp/rand31/p1192-park.pdf> (visited on Jan. 30, 2019).
- [PMS93] Stephen K. Park, Keith W. Miller, and Paul K. Stockmeyer. “Technical Correspondence: Response”. In: *Communications of the ACM* 36 (7 July 1993): *Special Issue on Computer Augmented Environments: Back to the Real World*, pp. 108–110. ISSN: 0001-0782. DOI: 10.1145/159544.376068. URL: <http://www.firstpr.com.au/dsp/rand31/p105-crawford.pdf#page=4> (visited on Jan. 30, 2019).